

Survival Model Evaluation: synthetic

Fitted with the survival-analysis-pipeline, on synthetic data with known ground truth

Source data	<code>data/strategies.csv</code> , drawn at seed 7
Rows	5,000 (93.2% with the ending observed, 6.8% censored)
Start dates	2021-07-01 to 2026-06-01
Evaluation	5 expanding-window temporal folds
Source of figures	<code>reports/synthetic/metrics.json</code> , regenerated by the command in section 9

1. Summary

This report evaluates two survival models fitted to synthetic data drawn at seed 7. It holds 5,000 rows, each observed from its start date. 93.2% have observed endings and 6.8% are censored, still running when observation stopped. Median observed duration, censored rows included, is 91 days, the scale on which every horizon and prediction below sits. The models predict, from what was on file at the start date, how long each row survives.

When evaluating the two models' performance, the apples-to-apples comparison is the fold-mean concordance, the share of pairs a ranking orders correctly where 0.500 is a coin flip. The pipeline uses walkforward validation. Each of the 5 temporal folds trains both models on one stretch of history and tests them on the next, and the 5 scores average. On concordance XGBoost AFT scores 0.781 and the Cox proportional hazards baseline 0.780. The two models effectively tie.

To attach an uncertainty range, which says how far the score could reasonably move, the AFT model is also scored with all 3,000 out-of-time test rows pooled into one list. That concordance is 0.781, with a 95% bootstrap interval of 0.772 to 0.789. Section 5 explains the two figures and how to read the interval.

Little of the pooled score is `asset_class` group membership. Ranking rows by their group's average prediction alone scores 0.538, and comparing only rows inside the same group scores 0.779, clear of the coin flip. Section 5 gives the decomposition.

An oracle ranking, defined in section 5, bounds every model at 0.820. Ranking by validation Sharpe scores 0.410, below the coin flip in every fold.

Both models are saved in one bundle, which records the boosted model as recommended for scoring new rows, on a near-tie margin.

All results are synthetic. The run validates the pipeline and asserts nothing about live data.

2. Motivation

Analyst notes:

An automated strategy search produces a queue of candidates that all clear validation, and they stop working at very different rates. How much capital a new strategy should receive, when its first review should be scheduled, and when it should be retired all depend on how long its edge will last. Treating every deployment the same misallocates in both directions. It overfunds strategies that will not survive the quarter and underfunds ones that would have run for a year.

The question this run poses, on a synthetic population built to mimic that setting, is whether discovery-time metadata carries enough signal to separate them. A strategy selected from a hundred thousand candidates carries more selection bias than one selected from a hundred, and a strategy whose walk-forward Sharpe was already sliding during validation had begun decaying before it was deployed. None of that is visible in a single validation statistic, and all of it is already recorded. On a real book, death would be a bookkeeping event, the date a strategy stops clearing the book's retention rule. The lifetimes a model learns there are partly a property of that rule. The generator abstracts this to a lifetime drawn at birth.

3. Data

Durations are measured in days.

Analyst notes:

Two generator settings matter for interpreting the results. 6% of rows are administratively retired independent of performance, which installs the censoring the evaluation must handle correctly. And lifetimes are drawn log-normally around what their drivers predict, at noise scale 0.55, the irreducible noise that keeps even the oracle in section 5 short of a perfect score.

The generator intentionally installs detectable structure in the data, so that recovering it is a test rather than an interpretation. Each of the six feature families shifts log survival time by a fixed amount, microstructure the most fragile and value-carry the most durable. Among the four asset classes, crypto is built to decay

fastest, a shift of -0.30 on log time against the fx-majors baseline. It is drawn for 15% of rows, so the class effect has to be recovered from a minority slice.

Censoring concentrates among recently discovered strategies, which have had the least time to fail before the observation cutoff, so the final fold's test block is far more censored than the population overall.

Figure 1 shows Kaplan-Meier survival curves by `asset_class`, the coarsest structure in the outcome before any model is fitted.




Figure 1. Kaplan-Meier survival curves by `asset_class`: the fraction of each group still running at each age, with censored rows counted for as long as they were observed. Groups beyond the seven most frequent, where present, are collapsed into (other) for this plot only.

4. Method

4.1 Model class

The pipeline fits two models on every run. The first is a boosted-tree accelerated failure time (AFT) model, XGBoost with its `survival:aft` objective, built for durations with incomplete observations. An observed ending tells the model the exact lifetime, and a censored row tells it only "at least this long", so every row contributes. It predicts each row's median survival time in days, and a log-normal curve around that median,

whose width is fitted once and shared by every row, gives the probability of surviving any horizon. The second is a Cox proportional hazards baseline, the standard linear survival model, fitted with lifelines' `CoxPHFitter` on the same features to answer whether the boosted model was necessary. Its coefficients are fitted under a penalty that pulls them toward no effect, which holds the fit steady when features move together but leaves the intervals around those coefficients approximate. The run recommends whichever scores the higher fold-mean concordance.

The two differ in what they assume. The AFT model assumes every row follows the same survival curve run on a faster or slower clock, so features change when things happen, never the curve's shape. The Cox model makes no assumption about that shape at all, estimating the curve from the data by counting who was still running at each age. In exchange it assumes each feature multiplies risk by the same factor at every age, which this report does not test. Which set of assumptions suits a dataset cannot be known in advance, which is why both are fitted and scored.

Numeric features pass through, missing values included (the Cox baseline gets train-window median imputation). Text columns are one-hot encoded with the vocabulary refit per training window, so a fold's features reflect only what was on file by its split date.

4.2 Temporal validation and label re-censoring

Evaluation uses 5 expanding-window folds ordered by start date: the earliest 40% of rows is burn-in, trained on, never tested, and each fold trains on every row started before its split date and tests on the next block (sizes in Table 2).

Training labels are re-censored at each split date. A row started long before a split may have died after it, and its label contains that future, so every post-split death is rewritten as a censoring at the split. Omitting this raises scores by importing the future, which makes it the most consequential detail in the pipeline. It is a standalone function (`temporal_folds.recensor`) with dedicated tests.

4.3 Selection and calibration

The boosted model's hyperparameters are selected once on the first fold's training window by an inner temporal split, the same past-then-future cut made inside that window. Concordance grades only whether rows are ordered correctly, and a model can order them well while drawing survival curves far too wide or too narrow. So selection is scored on held-out censored log-likelihood instead, which grades the whole predicted distribution against what was observed.

The predictive scale is the width of the boosted model's log-normal curve, one unitless number shared by every row. A larger scale spreads the model's probability over a wider range of lifetimes. Training used 0.6, a setting of the loss that shapes how the medians are fitted. Then, with the medians held fixed, the width that

best matches held-out outcomes on the training window's most recent stretch was measured at 0.64 and carried to the model refitted on the full window. The training value answers which curve width trains the best medians. The measured value answers which width matches the outcomes the trained medians actually got. The two need not agree. Every probability the boosted model reports here uses the measured value.

5. Results

5.1 Discrimination

Table 1. Concordance index by model, pooled over 3,000 test rows with 95% percentile bootstrap intervals over 500 resamples. Higher is better, and 0.500 is a coin flip. Fold-mean rows score each of the 5 folds separately and average. The interval resamples test rows with the fitted models held fixed, so it measures scoring precision, not stability across history. The fold spread is the guide to that.

METHOD	CONCORDANCE (HARRELL'S C)	95% INTERVAL
Oracle ranking on latent log-time (ceiling)	0.820	not resampled
XGBoost AFT (pooled)	0.781	0.772 to 0.789
XGBoost AFT (fold mean)	0.781	not resampled
Cox proportional hazards (fold mean)	0.780	not resampled
Rank by validation Sharpe	0.410	0.397 to 0.422

Each fold fits its own models. Because the AFT model predicts a survival time in days, and days is a universal unit that exists outside of an individual fold model, its predictions can be pooled into one list while Cox proportional hazards models cannot. A Cox model predicts a hazard risk score in relation to the average row in the window it was trained on, not universal time units. As a result, Cox scores from different folds cannot be compared in one list. For a fair comparison, the two models must be compared on fold mean concordance, which is each fold scored on its own and the 5 scores averaged.

The pooled row answers a different question. It scores all test rows in one list, which is enough data to attach an uncertainty range as seen in Table 1. The interval is the range the score stayed inside for 95% of resamples of the test rows. A narrow interval means the score is precise on this test set. And an interval that sits wholly above 0.500 means the model beats a coin flip even at its low end.

It is important to determine whether the model does more than recognize which group a row belongs to. In order to assess this, we split the pooled concordance by `asset_class`. Re-ranking rows by their group's average replaces every row in a group with the same score, so comparisons only ever run between groups. Evaluating the pooled concordance on those re-ranked rows gives a score of 53.8%. The boosted model scores

each row individually instead, which orders rows inside a group. Inside a group this ranking is correct 77.9% of the time, averaged over the 4 groups large enough to score (at least 50 rows and 10 observed endings), each weighted by its number of comparable pairs. If these values are greater than or equal to the original rankings, the model's performance is dominated by the within-group ranking.

The oracle ranking in Table 1 orders rows by the latent log-time the generator actually used. Nothing is fitted and the ranking predates the noise draw, so its 0.820 bounds every model. The gap to a perfect score is installed noise. A suite test asserts the model never outscores it, since beating perfect information indicates a leak.

Validation-Sharpe ranking scores 0.410 pooled, 0.354 to 0.427 across folds, never above 0.500. Reversed, it scores 0.590, its arithmetic complement, short of the model's 0.781.

Figure 2 plots the per-fold concordance from Table 2.

Table 2. Per-fold results. Censoring is the share of each fold's test rows whose ending was not observed.

FOLD	SPLIT DATE	TRAIN N	TEST N	CENSORED	AFT	COX	SHARPE	ORACLE
1	2023-06-22	1,999	600	2%	0.781	0.783	0.425	0.824
2	2024-02-03	2,600	600	2%	0.768	0.769	0.421	0.808
3	2024-09-01	3,199	600	2%	0.767	0.764	0.413	0.805
4	2025-04-03	3,800	600	9%	0.773	0.769	0.427	0.811
5	2025-11-03	4,399	600	39%	0.817	0.815	0.354	0.846




Figure 2. Concordance by fold. The boosted model spans 0.767 to 0.817. Folds differ in censoring mix, so cross-fold comparisons are indicative rather than exact.

5.2 Calibration

A majority of this report focuses on concordance. But concordance grades only whether models produce the correct ordering when comparing two rows. It cannot be used to assess the absolute accuracy of the predicted probabilities. To assess the absolute accuracy, we use the Brier score, the mean squared error of a probability forecast (lower is better). Censored rows can bias the Brier score, so we use the inverse probability of censoring weighting (IPCW) to compensate. At each horizon in Table 3, the models predict each row's chance of still running that length of time. The no-skill forecast predicts the same chance for every row, the share of the whole population still running at that horizon, estimated with the Kaplan-Meier method so censored rows count for as long as they were observed.

Table 3. IPCW Brier score by horizon. Lower is better, and the no-skill column is the score to beat, so a model above it adds nothing at that horizon.

HORIZON	AFT	COX	NO-SKILL FORECAST
90 days	0.149	0.151	0.249
180 days	0.132	0.133	0.182
365 days	0.050	0.051	0.056

Figure 3 plots predicted against observed survival at 180 days for both models. Points falling on both sides of the diagonal are noise. Points falling consistently on one side are bias in that direction.




Figure 3. Predicted against observed survival at 180 days for both models, each binned on its own predicted deciles. Observed frequencies are Kaplan-Meier estimates within each bin, so censored rows contribute correctly. The largest deviation is 0.046 in the boosted model's decile 10 and 0.035 in the Cox baseline's decile 8. Deviations are probabilities. Deciles are cut on predicted value and hold 300 to 300 rows each. In small or heavily censored deciles the observed value carries more uncertainty than three decimals suggest.

6. Feature analysis

6.1 The boosted model

The scores above grade how well each model ranks rows and how accurate its probabilities are, not which inputs it used. A feature attribution answers that. For one row and one feature, it is that feature's contribution to the difference between that row's prediction and the model's average prediction. The method used here is SHAP (SHapley Additive exPlanations).

Attributions sit on the log scale of survival time, so an attribution is a multiplier on predicted time rather than a number of days. An attribution of +0.3 multiplies that row's predicted survival time by about 1.35, and a negative one shortens it. That conversion follows from the scale and says nothing about this dataset.

Averaging a feature's attributions by size, ignoring sign, says how much that feature moves predictions across the explanation sample. A feature high in that average moves predictions a lot, and one low in it barely moves them.

Attributions are descriptive and in-sample, computed on the final boosted model's own training rows. Correlated features split credit by the model's internal choices as much as by the data, so directions are more trustworthy than magnitudes.

Figure 4 ranks the strongest features by that average, and Figure 5 shows, for each of those features, which direction it pushed and how much that varied from row to row.




Figure 4. The strongest features, ranked by the average size of their attributions across the explanation sample. Size is taken without regard to sign, so this says how much each feature moves predictions and not which way it moves them; the direction is in Figure 5. The scale is the same log scale of survival time, so a feature averaging 0.10 moves predicted survival time by about a tenth on a typical row. A yes/no flag from a text column is labelled by its value alone. Features below the strongest few are not drawn.




Figure 5. One row per feature, one dot per row of the explanation sample. The dot's position is the attribution that feature earned for that row, so dots left of zero shortened its predicted survival time and dots right of it lengthened. The dot's colour is that row's value of the feature, red high and blue low, so a row whose reds sit left of its blues is a feature whose high values shorten survival. How far the dots spread says how much the feature's effect varies from row to row, and a feature pinned at zero did nothing. A yes/no flag from a text column is labelled by its value alone and reads the same way, with red meaning the row carries that value.

6.2 The Cox baseline

The Cox baseline's drivers are its coefficients, reported as hazard ratios. A hazard ratio is the factor by which one unit of a feature multiplies the hazard and a hazard is the risk of an ending at a given age. The factor is the same at every age, so a ratio above 1 shortens survival. Figure 6 plots the strongest covariates.




Figure 6. The Cox Proportional Hazards model's strongest covariates with their 95% intervals on a log axis. Dots right of the line shorten survival and left lengthen it. The covariate whose effect the data pins down most firmly sits at the top, and only the strongest few are drawn. Ratios for `asset_class` compare against crypto (automatically chosen alphabetically) and excluded from the fit.

7. Interpretation

Analyst notes:

Why ranking by validation Sharpe inverts. The generator installs this inversion deliberately, and the design mimics how real-life strategy decay behaves. Every strategy enters the population by clearing a validation-Sharpe threshold. In the case of our synthetic study that threshold is 0.8. The generator draws an observed Sharpe as the sum of true edge, an overfitting component, and noise, which is the standard winner's-curse setup. Conditioning on that sum exceeding a bar means the survivors with the highest scores are disproportionately the ones whose overfitting and noise broke upward by chance. Past the bar, additional observed Sharpe is more likely inflation than edge. The inflated strategies are the overfit ones, which decay fastest. So higher validation Sharpe actively predicts shorter working life. It scores 0.410 pooled and sits below a coin flip in every fold. The inversion is a property the generator's design implies, and the run confirms the pipeline detects it. The winner's curse is a general statistical effect, so this likely generalizes to real strategy

populations. But the data here is purely synthetic, and the run is evidence about the pipeline, not about any real book. This is also what makes the modeling problem worth posing. If the backtest metric ranked survival correctly there would be nothing to add. A meta-model earns its place precisely because selection has already consumed the obvious signal.

Reading the tie. The boosted model and the Cox baseline land at 0.781 and 0.780 on the fold mean, a gap inside the bootstrap interval, which I read as a tie. The synthetic data generator's observable structure is close to additive, which leaves little for trees to find beyond what a penalized linear model in the log-hazard captures. I report the tie rather than adding interactions to the generator until the headline model wins, because a benchmark tuned until it loses is not a benchmark.

The calibration failure worth keeping. An earlier version of this pipeline selected hyperparameters by concordance and looked correct, then lost to a no-skill forecast on 365-day Brier score. Concordance is invariant to the predictive scale, so the search had settled on a distribution width far too wide and nothing in the training loop objected. That failure is why selection now uses censored log-likelihood and why every report this pipeline produces grades probability quality separately from ranking.

What the attribution shows, checked against the mechanism. The walk-forward statistics that dominate the ranking are not inputs to survival time in the generator. They exist only as noisy proxies for the latent edge-versus-overfit split that is the actual driver, and the model found them and leaned on them. Feature-family and asset-class effects return with the installed signs, among them the crypto shift of -0.30 described in the data note. High walk-forward positive fraction pushes predicted survival up, and a steeply negative walk-forward Sharpe slope pushes it down hard. High fold-to-fold Sharpe dispersion is penalized. All three directions match the generative design, where consistency rises with true edge and falls with overfitting. Validation Sharpe itself is nearly absent from the attribution ranking, which is the winner's curse seen from inside the model. The supportable reading of the magnitudes is that the model uses positive fraction about twice as much as Sharpe dispersion, not that positive fraction is twice as important, since the three walk-forward statistics are correlated proxies of one latent quantity.

Intended use, refit on a real book. Refit on production metadata, the model would be a capital-allocation and review-scheduling prior, never a trade signal. Predicted median lifetime sets the first review date and is one input to initial sizing, and the full curve gives a decay schedule to compare live performance against. It prices day-one information only. Once a strategy is live, realized performance carries information no discovery-time metadata can. Before any such use, the same temporal cross-validation with the same re-censoring has to be repeated on production metadata, and the resulting concordance should be expected to come in lower, with wider intervals.

8. Limitations

Censoring may be informative. The evaluation assumes rows stop being observed for reasons unrelated to their risk. Early exits of rows about to end anyway bias survival estimates upward.

Analyst notes:

The generator encodes a prior, not evidence, so the concordance here is an upper bound on production performance. Lifetimes are drawn and resampled independently, understating real-book uncertainty, and administrative retirement is modeled as independent censoring, with competing risks left as future work. Every number here comes from one generator seed, so data variance and model variance cannot be separated.

The boosted model gives every row one curve shape. The predictive scale is a single fitted number, so the model runs a row's clock faster or slower but never changes the curve's shape. Per-row width is future work.

Harrell's concordance is biased under heavy censoring, usually upward. It scores only the pairs censoring leaves comparable, which under heavy censoring over-represent short observed lifetimes. The most censored test window is 39% censored, so read its rows in Table 2 with the most doubt.

9. Reproducing this run

Python 3.11 or later. From the repository root:

```
pip install -r requirements.txt
python -m pytest
python scripts/run_synthetic_pipeline.py
```

Every number is read from the run's `metrics.json` at build time, and a figure the prose never cites fails the build. `requirements.txt` pins the exact versions the committed numbers came from.

Generated from `reports/synthetic/metrics.json` by `scripts/run_build_report.py`. 5,000 rows drawn at seed 7, 3,000 out-of-time test rows.